

AI PRD

AtliQ Sales & CRM Memory Assistant · branded in-product as “AI Sales Command Center”

Learner Name	Aishwarya Odhala
Project Name	AtliQ Sales & CRM Memory Assistant (Capstone 1)

Field	Detail
Product Name	AtliQ Sales & CRM Memory Assistant
Date	25 / 07 / 2026
Version	V1.0 — Draft
Status	Draft
Problem Statement	Founder-led sales teams at AtliQ lose revenue not from a lack of leads but from a lack of memory — critical signals live in emails and meeting notes while the CRM silently goes stale, and nobody has time to reconcile the two.

PART A — Traditional PRD Components

T1. Executive Summary

Field	Your Answer
What are we building?	An AI-powered memory layer that reads AtliQ's CRM, emails, and meeting notes together and surfaces mismatches, missing information, time-sensitive follow-ups, and dropped commitments — with every finding traceable to its source and every suggested action requiring human approval. It also prioritizes leads two ways (overall, and by issue type), highlights deals that quietly deserve a second look, and remembers what's changed since the last check.
Why does it matter?	Deals worth tens of thousands of dollars (e.g. a \$95,000 account) are going stale purely because the CRM doesn't reflect reality — this prevents that revenue leak without adding manual admin work for already-stretched founders.
Who is it for?	Primary: AtliQ's deal owners (Bhavin, Dhaval, Karandeep) and sales executive (Jay). Secondary: prospects/clients, who benefit indirectly through faster, better-informed follow-up.
How will we know it worked?	A measurable drop in high-value leads that sit untouched for 30+ days despite real, ongoing client engagement.
Target launch	Prototype complete now (capstone submission); pilot with AtliQ's own sales team targeted for Q4 2026.

❀ *A \$95,000 deal shouldn't go quiet just because nobody had time to reopen an email thread.* ❀

T2. Problem Statement & Background

Field	Your Answer
Problem statement	AtliQ's small sales team — two founders, the CEO, and one dedicated salesperson — juggle closing deals with running the company, leaving no time to keep the CRM honest — so the real state of every relationship lives scattered across emails and meeting notes instead of the system meant to track it.
Current state	A seller has a call or exchanges emails with a prospect; they may, or may not, log it in the CRM days later, from memory — often missing details like next steps, exact commitments, or a changed deadline. No process exists to reconcile what the CRM says against what actually happened.
Pain points	These are the friction points a deal owner personally feels while doing the job — each grounded in something we actually found in the dataset: • “I don't always know a deal has genuinely moved forward until I happen to reopen the email thread” — FinEdge Bank's compliance clearance sat unnoticed for weeks while the CRM still read “Proposal Sent.” • “I lose the ability to follow up the moment one key detail never gets logged” — BlueOrchid Hotels' real pricing conversation had no CRM contact email captured. • “I don't always know a colleague is already talking to the same company” — Meridian Healthcare was entered twice, by two different owners, without either noticing. • “I make promises in the moment with nothing reminding me to keep them” — a commitment to send case studies “by Friday” (ByteBridge) was never fulfilled or flagged. • “Some relationships never even make it onto my radar as an open lead” — real prospects (Harbor & Finch Advisory, Lumino Fitness, PixelWorks Agency) were never added to the CRM at all.
Business impact	Stalled high-value deals directly cost revenue; duplicated outreach risks looking unprofessional to prospects; dropped commitments erode trust before a deal even closes.
Why now?	AtliQ's inbound pipeline (referrals, conferences, LinkedIn) is scaling faster than its founder-led sales process can absorb — the gap between leads coming in and the system's ability to track them will only widen without intervention.

T3. User Personas & User Stories

T3A. User Personas

Persona	Description
Primary user	Bhavin (representing AtliQ's deal owners generally) — co-founder splitting time between running AtliQ and personally closing deals; owns several of the highest-value, most time-sensitive leads in the pipeline; needs to trust the CRM reflects reality without spending time manually maintaining it.
Secondary user / stakeholder	Jay (sales executive) — handles higher lead volume day-to-day, same underlying pain. Karandeep (CEO) — needs pipeline visibility without personally digging through every deal.

T3B. User Stories

As a...	I want to...	So that...	Acceptance Criteria
---------	--------------	------------	---------------------

Deal owner	see which leads have a mismatch between the CRM and what actually happened	I don't let a genuinely progressing deal go cold	Given a lead has an email/note contradicting its CRM status, when I open the dashboard, then that lead appears flagged with the specific contradiction shown
Deal owner	see missing information that's blocking follow-up (like a contact email)	I can fix the gap instead of losing the lead entirely	Given a lead has a linked document but a missing key CRM field, when I view that category, then the lead appears with the specific missing field named
Deal owner	have soft timing language ("check back after the board meets") converted into a real urgency signal	I know what's due now versus months away	Given a document mentions a future event tied to follow-up, when that event's date has passed or is imminent, then the lead is flagged due-now
Deal owner	get a suggested next action drafted for my most urgent leads	I can act quickly without composing a response from scratch	Given a lead is scored Urgent, when I open its detail view, then a suggested action is shown as a draft, clearly labelled as not sent
Deal owner	see the exact original email or note behind any AI finding	I can verify it myself before acting	Given a finding is shown, when I click "linked documents", then the original source text is displayed read-only

T4. Functional Requirements

ID	Requirement	Priority	Status
FR-1	The system shall load the CRM export, all email threads, and all meeting notes, and link them by company name, contact email, and lead ID.	Must	Open
FR-2	The system shall detect and display mismatches between CRM status/fields and linked emails/notes, with evidence quoted verbatim.	Must	Open
FR-3	The system shall flag CRM records missing key fields (contact email, next follow-up date, owner) where linked documents show real engagement.	Must	Open
FR-4	The system shall convert soft timing language into a due-now vs. future urgency classification relative to the current date.	Must	Open
FR-5	The system shall identify unanswered client questions and unfulfilled AtliQ commitments from email/note text.	Must	Open
FR-6	The system shall detect likely duplicate CRM records for the same company/contact.	Should	Open
FR-7	The system shall identify companies mentioned in emails/notes with no corresponding CRM record.	Should	Open
FR-8	The system shall generate a suggested next action (draft only) for leads scored Urgent, without sending anything automatically.	Should	Open
FR-9	The system shall reject any AI-generated finding that cannot be matched to a verbatim quote in its source document.	Must	Open

FR-10	The system shall allow filtering the dashboard by lead owner (Bhavin, Dhaval, Karandeep, Jay), including an explicit 'Unassigned' group for leads with no owner set.	Must	Open
FR-11	The system shall generate a brief AI-written executive summary synthesizing overall pipeline state, once per analysis run.	Must	Open
FR-12	The system shall compute an evidence-completeness score per lead from real data-completeness signals (linked document count, key CRM fields present, contact recency) — never a fabricated AI confidence value.	Must	Open
FR-13	The system shall rank leads within each category by relevance to that specific category, not by overall urgency, so different categories surface different top leads.	Must	Open
FR-14	The system shall identify Hidden Opportunity leads — medium urgency but showing strong underlying signals such as confirmed budget or an engaged decision-maker — as a distinct section.	Should	Open
FR-15	The system shall compute a qualitative Impact indicator (High / Medium / Low) per Urgent lead from deal value, urgency, and staleness.	Should	Open
FR-16	The system shall maintain a bounded history of past analysis runs and report what's new, resolved, and persisting since the last check.	Should	Open
FR-17	The system shall display each lead's original source documents (emails/notes) read-only, alongside any AI finding that cites them.	Must	Open
FR-18	The system shall display each lead's source channel (Referral, LinkedIn, Conference, Website, Cold Outreach) as recorded in the CRM export.	Could	Open
FR-19	The system shall report total analysis time and lead count after each run, to make processing efficiency visible.	Could	Open

T5. Non-Functional Requirements

Category	Requirement & Target
Performance	Full analysis of ~40–45 leads completes in a few minutes (async batch); on-demand actions (e.g. generate email) return within ~10s.
Reliability / Uptime	Prototype-grade; cached results keep the dashboard usable even if the AI API is briefly unavailable, since findings persist locally.
Security	API key stored server-side only (.env), never exposed client-side; no data sent beyond the model provider.
Privacy & Compliance	All prototype data is synthetic; a production version would follow AtliQ's data residency/retention policy and exclude any field not needed for the task.
Scalability	Scoped to AtliQ's current pipeline size (~40–50 leads); caching by content hash avoids re-analysing unchanged leads as volume grows.
Integrations	Reads CRM data via file export in the prototype; production would connect read-only to AtliQ's live CRM and inbox.

T6. Out of Scope

Out of Scope Item	Reason / When It May Be Revisited
Automatic write-back to the CRM	Stays human-owned by design (core guardrail); revisit only after a long track record earns that trust.
Sending emails/messages to clients automatically	Same reasoning — this is a decision-support tool, not an outreach automation tool.
Real-time/live inbox monitoring	This is a periodic check-in tool (run on demand), not a continuously running background agent — a separate, later capability.
Full conversational “AI Sales Coach” chat interface	Identified as valuable; deferred to keep prototype scope on the core memory-layer capability. Natural v2 addition.
Gmail / WhatsApp integration	Raised directly by the client at kickoff as a future direction; explicitly out of scope for this prototype, which reads from file exports only.

T7. Timeline & Milestones

Milestone	Target Date	Owner / Notes
PRD approved	Jul 2026	PM (self)
Prototype complete	Jul 2026	PM + Claude Code
Internal alpha (capstone)	Jul 2026	PM
Stakeholder demo	Aug 2026	PM
Pilot with real AtliQ team	Q4 2026	Sales team + PM

PART B — AI-Specific PRD Components

1. Problem & User — JTBD Framing [AI-SPECIFIC]

Field	Your Answer
Job To Be Done	When I sit down to check my pipeline, or before a client call, help me see every unresolved signal about that relationship — across CRM, email, and notes — in one place, so I can follow up with confidence instead of reconstructing the deal's history from memory.
Trigger	A deal owner opens the tool to plan their week, or before contacting a specific client.
Outcome	A ranked, evidence-backed list of what's wrong, missing, or time-sensitive per lead, plus a draft suggested action for the most urgent ones — reducing effort from re-reading every email to reviewing one short summary.
Why AI is justified	Manually cross-referencing 40+ CRM records against 50 emails and meeting notes every week isn't realistic for a small team who are also running the company. The

	volume and unstructured nature of the text make this a strong fit for language-model extraction — provided every output stays grounded in verifiable source text.
Secondary stakeholders	Sales leadership (Karandeep) needing pipeline visibility without digging through every deal; indirectly, clients who benefit from more timely, informed follow-up.

✿ *The CRM isn't broken — it just only knows what someone remembered to tell it.* ✿

2. Opportunity & AI Capability [AI-SPECIFIC]

Opportunity	AI Capability Used
Detecting mismatches between structured CRM data and unstructured email/note text	Classification + Structured Output
Extracting missing/incomplete information signals	Structured Output
Converting informal timing language into an urgency classification	Classification + light Generation
Drafting a suggested next-action email	Generation (grounded, template-constrained)

3. Model Strategy — Build vs Buy vs Fine-tune [AI-SPECIFIC]

Decision	Your Answer
Approach	Prompt-only with schema-constrained structured output. Not RAG, not fine-tuned, not agentic — each lead's linked-document set is small enough to pass directly in context.
Why this approach	- Per-lead document volume is small (a handful of emails/notes), so full-context prompting beats building a retrieval pipeline for marginal benefit. - A defined JSON schema per finding category makes outputs directly renderable and easy to grade for correctness.
Why not fine-tune in v1	No labelled training data exists yet for “correct mismatch extraction” at the volume needed to fine-tune responsibly; prompt + structured output is the right default, revisited once real usage data (human agreement/override rates) accumulates.
Build vs Buy	Built as a custom layer (Streamlit + Claude API) rather than buying an existing sales-intelligence tool, because the core value is bespoke reconciliation logic tailored to AtliQ's specific data and guardrails — not an off-the-shelf capability.
Model requirements	Structured/schema-constrained output support; strong instruction-following to enforce the “quote verbatim or don't claim it” grounding rule; low hallucination rate on extraction tasks.
Latency target	Async/batch acceptable for the main analysis (a few minutes is fine — periodic check-in, not live chat); on-demand actions (e.g. generate email) under ~10s.

4. Context & Data Pipeline [AI-SPECIFIC]

4A. What Goes In

Input	Source & Notes
CRM export (crm_export.csv)	AtliQ's CRM system export, file-based in the prototype. Refreshed on-demand ("Run/Refresh analysis"); production would sync more frequently.
Email threads	Client/prospect correspondence, provided as text per lead. Same refresh cadence; production would connect read-only to the sales inbox.
Meeting notes	Call/meeting summaries logged by sellers. Same refresh cadence as above.
Prior analysis history	The tool's own cached findings from previous runs, used only to compute what's new / resolved / persisting — not re-sent to the model as fresh context.

4B. Deliberately Excluded From Context

Excluded Attribute	Reason for Exclusion
AtliQ-internal financial data beyond a deal's est_value_usd	Not relevant to relationship-tracking; unnecessary exposure of sensitive business data to the model.
Personal/protected attributes of any client contact	Must not influence any AI output — bias and privacy risk with no legitimate task justification.
Data from other AtliQ business lines (e.g. delivery/project execution)	Out of scope per the problem definition; keeps the model focused and reduces irrelevant noise.

5. Behaviour & Output Specification [AI-SPECIFIC]

5A. Standard Output

Output Field	Description
Primary output	A categorised list of findings per lead (mismatch / missing info / follow-up-timing / open question / deprioritise signal), each a short plain-language description, plus an overall priority score, a qualitative Impact indicator for Urgent leads, and, where applicable, a suggested next action. A pipeline-wide executive summary sits above all of this.
Supporting evidence	Required — every finding must cite a verbatim quote from its source document; findings that can't be matched to real source text are automatically dropped before display.
Confidence signal	Framed as "evidence completeness," not a fabricated confidence percentage — computed from real signals (linked document count, whether key CRM fields are populated, recency of last contact), since the model doesn't output a genuine calibrated confidence value.
Format	Structured JSON per lead, rendered as categorised cards/lists in the Streamlit dashboard.

5B. Required Behaviour States


State	What Triggers It	Required Behaviour
High confidence	Finding passes the verbatim-quote grounding check	Shown with its source evidence and category tag; still requires human review before any related action is taken
Low confidence	Thin evidence — few linked documents, stale contact	Reflected via a lower evidence-completeness score with the specific gaps listed; shown, not hidden

No data / insufficient info	Lead has no linked documents	Shown as “not yet analysed” or with no findings — never a fabricated finding
Below-threshold output	Finding fails the verbatim-quote grounding check	Silently dropped from what's displayed; the debug panel reports how many were dropped, for builder transparency
Suspected manipulation / injection	Source document contains embedded instructions	Documents are treated strictly as data to extract from, never as instructions; the extraction prompt constrains the model's role

5C. User Flow — AI Decision Flow

```

flowchart TD
    Start([User clicks Run/Refresh]) --> LoadData[Load CRM export, emails, meeting notes]
    LoadData --> LoadCheck{Data loaded successfully?}
    LoadCheck -->|No| LoadError[Show load error, stop]
    LoadCheck -->|Yes| Link[Link documents to leads]
    Link --> Structural[Structural checks: duplicates, not-in-CRM]
    Structural --> CacheHit{Cached result exists?}
    CacheHit -->|Yes| UseCache[Use cached finding]
    CacheHit -->|No| CallAI[Call Claude: extract findings]
    CallAI --> AIOk{Call succeeded?}
    AIOk -->|No, retry| CallAI
    AIOk -->|Still failing| Queue[Mark not-yet-analysed, add to review queue]
    AIOk -->|Yes| Ground{Verbatim quote match?}
    Ground -->|Yes| Keep[Keep finding]
    Ground -->|No| Drop[Drop finding, log for debug]
    Keep --> Score[Score urgency + category-specific ranking]
    Drop --> Score
    Score --> Trail[Compare to last run: new / resolved / persisting]
    Trail --> Summary[Generate executive summary]
    Summary --> Evidence[Compute evidence-completeness per lead]
    Evidence --> Display[Display: Weekly Attention, categories, Effort vs Payoff, Hidden Opportunity]
    Display --> Queue
    Queue --> Display
    Display --> HumanGate[Human reviews findings + linked source documents]
    HumanGate --> Filter[Optional: filter by owner]
    Filter --> GenEmail[Optional: generate draft follow-up email]
    GenEmail --> End([Human copies, edits, sends manually])
    Filter --> End
  
```

[Paste your MermaidFlow diagram image here]

6. Guardrails & Failure Modes [AI-SPECIFIC]

Failure Mode	What Goes Wrong	Required Guardrail	Owner
Hallucination	Model claims a finding not actually in the source	Every finding must match a verbatim quote from its linked document; failures are auto-dropped before reaching the user	PM/Eng
Bias amplification	Scoring could favour leads based on non-business signals	Only business-relevant fields (deal value, staleness, evidence) drive scoring; protected attributes excluded from context entirely	PM/Eng
Automated action without human	AI could send an email or edit the CRM directly	System is read-only to the CRM/inbox; all suggested actions are drafts requiring explicit human approval, by architecture not just policy	PM/Eng
Prompt injection	A client email could contain manipulative embedded text	Source documents are always data to extract from, never instructions; extraction prompt explicitly constrains the model's role	Eng
Explainability gap	A reviewer asks why a finding was flagged	Every finding links to its source document, viewable read-only via the linked-documents/source viewer	PM/Eng

✿ *Every suggestion is grounded in real evidence, and every action still needs a human yes.* ✿

7. Eval Criteria — Quality Bar Before Shipping [AI-SPECIFIC]

Eval Dimension	Definition & Threshold
Functional eval	Correct = the finding's stated issue is genuinely present in the source document on human review. Target: 100% of five validated dataset examples (FinEdge, BlueOrchid, Meridian, Trustline, ByteBridge) correctly surfaced.
Safety eval	Zero tolerance for ungrounded findings reaching the display layer — enforced automatically by the grounding check, not manual review.
Content quality gate	≥80% of a sampled run's findings must be usable/accurate without major correction.

Human agreement rate	PM (as domain reviewer) agrees with $\geq 80\%$ of flagged findings on a representative sample.
Launch gate (single condition)	All five validated real-data examples must be correctly, specifically surfaced by the running prototype — non-negotiable before any demo.

8. Success Metrics [AI-SPECIFIC]

Layer	Metric	Target / Threshold
North Star	Count of high-value leads (est_value_usd above ~\$50k) with 30+ days of no CRM activity despite real engagement in linked documents. This one metric directly maps to all three focus areas raised at client kickoff: fewer leads lost to unstructured data, shorter follow-up delays, and less manual analysis time — a stale high-value lead is the visible symptom of all three failing at once.	Trend toward zero for actively engaged leads over the pilot period
Output (shallow)	Number of findings surfaced per run; number of suggested actions generated	Directional only, tracked not targeted
Trust signal	Human agreement rate with flagged findings; rate suggested actions are used as-is vs. edited vs. discarded	$\geq 80\%$ agreement
Guardrail metric	Rate of findings dropped by the grounding check; count of automated CRM writes or sent emails	Grounding-drop rate tracked (healthy, not minimised); automated writes/sends must stay at 0

✿ *North Star: watch the stale, high-value leads shrink toward zero — that's the whole point.* ✿

9. Scope & MVP Boundary [AI-SPECIFIC]

v1 — Ships	Later — Explicitly Deferred
- Mismatch, missing-info, follow-up/open-question, and deprioritise-signal detection, grounded in verbatim quotes - Duplicate and not-in-CRM detection - Weekly Attention carousel + six category views, each ranked by relevance to that specific category - Effort vs. Payoff chart and Hidden Opportunity section - Evidence completeness score, linked-document source viewer, memory trail (“since your last check”) - AI-generated executive summary and qualitative Impact indicator (High/Medium/Low) - Draft suggested actions (next step + follow-up email), human-approved only	- Full conversational “AI Sales Coach” chat interface — deferred to keep scope on the core memory-layer capability - Direct live CRM/inbox integration — prototype uses file exports - Automated actions without human approval — only after a real trust track record - Deal-closing probability prediction — needs historical outcome data not in this dataset

Owner-based filtering (including an explicit Unassigned group), lead-source display, and a run-time/efficiency indicator	
--	--

10. Rollout & Release Strategy [AI-SPECIFIC]

Stage	Definition & Gate Condition
Shadow mode	Runs alongside the existing manual process with no changes made on its output yet. Gate: PM manually validates a sample against real emails/notes, confirms $\geq 80\%$ agreement.
Closed beta	Rolled out to AtliQ's own deal owners only. Gate: structured weekly feedback confirms the tool is actually checked and found useful.
Phased rollout	Expand to a second internal team or broader lead volume as AtliQ's pipeline grows. Gate: guardrail metrics hold steady at scale.
Full release	Gate: sustained usage and positive trust signals over multiple consecutive weeks.
Kill switch conditions	Any confirmed fabricated (non-grounded) finding reaching a user despite the guardrail; any accidental automated CRM write or sent email — either triggers immediate rollback and guardrail review.

11. Open Questions & Assumptions Log [AI-SPECIFIC]

The four rows below were raised directly by the client at kickoff and are addressed here so the resolution lives with the rest of the product decisions, not just in a call recording.

Open Question / Assumption	Ticket Type	Owner	How We Resolve It
How should unassigned leads be handled? (client kickoff)	Resolved	PM	Leads with no owner are shown under an explicit 'Unassigned' group (FR-10) rather than being silently excluded from any filtered view.
What should be the final North Star metric? (client kickoff)	Resolved	PM	Stale high-value leads trending to zero — see Section 8. Chosen because it's the one metric that reflects all three of the client's suggested focus areas at once.
Can follow-up dates be auto-generated from existing data? (client kickoff)	Partially resolved	PM + Eng	Soft timing language is already converted into a due-now/future signal (Follow-up & Open Questions category). Generating an exact calendar date automatically is a reasonable v1.1 enhancement, not yet built.
Is Gmail/WhatsApp integration in scope for the prototype? (client kickoff)	Resolved	PM	No — explicitly future scope (see Section T6, Out of Scope). The prototype reads from file exports only.
Is extraction reliably accurate across a much larger, messier real dataset?	Capability spike	PM + Eng	Run on a larger labelled sample once real AtliQ data is available; measure precision against human-reviewed ground truth
Will founders actually trust and act on AI-suggested next actions?	Assumption to test	PM	Shadow/beta mode; compare to what a human would have written; gather direct feedback

Should the tool connect live to the CRM/inbox instead of file exports?	Infra spike	Engineering	Prototype the integration once file-based value is proven; evaluate refresh-frequency needs first
Is Groq (used during prototyping for cost) or Claude the right production default?	Capability spike	PM + Eng	Compare cost vs. quality on a larger sample before committing

PRD Pressure Test — Before You Present

1. What does the AI do when it is NOT confident? → See Section 5B: it's shown with a lower evidence-completeness score and the specific gaps named, never hidden or fabricated.
2. What data is deliberately excluded from context, and why? → See Section 4B: protected personal attributes and unrelated business data, to avoid bias risk and irrelevant noise.
3. What is the single condition that must be true before this ships? → See Section 7: all five validated real-data examples must be correctly surfaced.
4. Did you write a JTBD or a feature description in Section 1? → A JTBD — framed around the trigger (checking the pipeline), the outcome (a ranked, evidence-backed summary), and the benefit (following up with confidence).
5. Would an engineer know what to build from this document alone? → Yes — Sections 4, 5, and the Mermaid flow in 5C define inputs, outputs, behaviour states, and the end-to-end decision flow explicitly.