

AtliQ Sales & CRM Memory Assistant

Cost Estimation — Summary Dashboard

Monthly AI tool cost	Annual AI tool cost	Monthly savings vs. manual	% cost saved	Annual savings
\$53.24	\$638.87	\$446.38	89.3%	\$5,356.52

ROI — for every \$1 spent on the AI tool, this saves:

9.4x

in founder + sales rep time value

WHY THIS APPROACH, AND HOW CACHING KEEPS IT CHEAP

- Prompt-based, not RAG or fine-tuned: each lead's linked emails/notes are small enough to pass in full to Claude Sonnet 5 — simpler and more reliable than a retrieval pipeline at this data volume, and there's no labelled dataset yet to fine-tune on responsibly.
- Without caching, every Refresh would re-analyze the entire pipeline from scratch — at \$0.0698/lead × 43–50 leads, that's roughly \$3 every single click, every day.
- With caching, a lead is only ever re-analyzed when its linked emails/notes genuinely change, or when the extraction prompt itself changes (a PROMPT_VERSION bump). An unchanged lead costs \$0 on every repeat Refresh.
- This is exactly why real spend can spike to \$3–4 right after a prompt/schema update — the whole cache invalidates once, and every lead needs one fresh pass — then drops back to near-\$0 per click once the prompt is stable again. That spike is a one-time development cost, not a recurring production cost.

INFERENCE — MODEL INFERENCE COST, GROUNDED IN REAL MEASURED USAGE

Yellow-highlighted rows in the Excel version are adjustable assumptions.

Model & Pricing

Model	claude-sonnet-5
Input cost (\$ / 1M tokens)	\$3.00
Output cost (\$ / 1M tokens)	\$15.00

Per-Lead Analysis Cost — Measured

Full-run cost (43 leads, 6 categories, actual billing)	\$3.00
Average cost per lead (measured)	\$0.0698

Monthly Volume Assumptions — Steady State

Total leads in pipeline	50
Leads needing fresh analysis / week	15
Run/Refresh clicks / month (updated: ~1x per day)	30
Follow-up emails generated / month	25

Cache-Only Technical Floor — If Nobody Ever Force-Refreshes (Rarely Realistic Alone)

Lead analysis cost / month	\$4.53
Executive summary cost / month	\$0.21
Follow-up email cost / month	\$0.16
Portfolio Insights charts (5 of 6, \$0)	\$0.00

Total monthly cost (technical floor)	\$4.90
Total annual cost (technical floor)	\$58.85

This floor assumes every Refresh is a cache hit unless a lead is genuinely brand-new — in practice, teams don't rely on cache alone; see 'Standard Company Usage' below for the realistic number.

What a Prompt / Schema Change Costs — A One-Time Event

Leads re-analyzed on a PROMPT_VERSION change	50
Cost of one full re-analysis after a schema/prompt change	\$3.49

This is what you were actually seeing this week — each prompt fix (mismatch dedup, competitor-mention rule, pricing-pushback threshold) bumped PROMPT_VERSION and forced one full re-analysis on the next Refresh. It is not a recurring per-click cost — once the prompt is stable, Refresh returns to near-\$0 unless leads genuinely change.

Routine Force-Refresh — Standard Company Practice (3–4x/Week)

A real team doesn't rely on cache alone — checking with a forced fresh pass several times a week is the realistic, expected pattern, not an edge case. This is the primary planning number, not a caution footnote.

Force-refresh uses / week	3.5
Leads re-analyzed per force-refresh	50
Cost per force-refresh	\$3.49
Weekly cost (force-refresh only)	\$12.21
Monthly cost (force-refresh only)	\$52.87
Yearly cost (force-refresh only)	\$634.40

Total Monthly Cost — Standard Company Usage (Recommended Headline)

Daily normal Refresh + routine force-refresh 3–4x/week, combined. This is the number to quote when someone asks 'what will this actually cost us.'

Force-refresh cost / month	\$52.87
Executive summary cost / month	\$0.21
Follow-up email cost / month	\$0.16
Total monthly cost — standard usage	\$53.24
Weekly	\$12.30
Yearly	\$638.87

COST SAVINGS — AI TOOL VS. MANUAL REVIEW

What this replaces: a founder or Jay manually re-reading emails and notes against the CRM.

Human Hourly Rate — Market-Sourced

Average CEO / Founder salary (\$/year)	\$150,000
Average Sales Executive salary (\$/year)	\$90,000
Standard working hours / year	2,080
CEO / Founder — hourly rate	\$72.12
Sales Executive — hourly rate	\$43.27
Blended hourly rate (average of both)	\$57.69

Seed-stage CEO average ≈ \$147K (Kruze Consulting, 2025); Sales Executive ≈ \$81K–\$93K (ZipRecruiter / Indeed, 2025–26). Both rounded. This blended rate replaces the earlier flat \$90/hr placeholder — and it's lower, so prior savings figures were, if anything, conservative.

Manual Effort Baseline

Manual review time / lead needing attention (minutes)	8
Leads needing fresh review / week	15
Manual review hours / month	8.66
Monthly human-equivalent cost	\$499.62

Savings — Standard Company Usage (Recommended Headline)

AI tool cost / month	\$53.24
For comparison: cache-only technical floor	\$4.90
Savings / month	\$446.38
% savings	89.3%
Annual savings	\$5,356.52
ROI — return per \$1 spent	9.4x

Break-Even

Break-even manual review time per lead (minutes)	0.85
--	------

At any manual review time above this many minutes/lead, the AI tool is cheaper. The assumed 8 min/lead is well above this line.

MODEL CHOICE, INFRASTRUCTURE & KEY ASSUMPTIONS

Which model?	Claude Sonnet 5 (claude-sonnet-5), via the Anthropic API.
Why this model?	Strong instruction-following for the grounding/verbatim-quote guardrail; native structured/schema-constrained output support; more than sufficient quality for extraction and comparison tasks at this scope.
Why not RAG / embeddings?	Each lead's linked-document set is small (a handful of emails/notes), so passing the full set directly in context is simpler and more reliable than building a retrieval pipeline at this data volume.
Why not fine-tuning?	No labelled training data exists yet at the volume needed to fine-tune responsibly. Prompt + structured output is the right default for v1.
Hosting	Streamlit Community Cloud — free tier, \$0/month.
Code hosting	GitHub — free, public repository.
Storage / caching	Local JSON files (analysis_cache.json, findings_history.json) — no database cost at this scale. A lead's cached result is invalidated only by new linked-document content or a PROMPT_VERSION change.

This model is fully formula-driven in its Excel form — every number traces back to the Inference or Cost Savings sheet.