

AI Sales Command Center

An AI Memory Layer that keeps your CRM aligned with every customer conversation.

DATA SOURCES



Emails
34 Threads



Meeting Notes
16 Notes



CRM Export (.csv)
40 Leads

AI MEMORY LAYER



Designed & Developed by

Aishwarya Odhala

Business Analyst | AI Product Builder

AI Sales Command Center

OverviewAll ownersBhavinDhavalJayKarandeepUnassigned

Executive Summary

Urgent Attention
12
Leads need action

Missing Information
8
Leads missing key info

Follow-ups Due
15
Open questions / commitments

New Since Last Check
7
Clients with updates

Consider Deprioritizing
6
Signs of cooling interest

This Week's Attention

1234

100 NovaPharma (L-1029) **URGENT**

Mismatch: Compliance cleared but CRM still "Proposal Sent"

Deal Value
\$95,000

Evidence completeness
85 / 100

Suggested action: Follow up with Pranav to start the vendor security questionnaire (78 questions on hosting/access control).

View Details →



AtliQ

Agenda

Five stops, one story.

1

The Problem

Why sales teams lose deals — and what's really going on underneath

2

The Investigation

Real evidence, the root cause, and the decision it led to

3

The Design

How the AI was scoped, guarded, and chosen on purpose

4

The Proof

Real cost, real metrics, and where this approach hits its limits

5

The Product

A feature tour — then the live demo


A QUESTION WORTH SITTING WITH

Why Do Sales Teams Lose Deals?

Ask most people, and you'll get the wrong answer.

IT'S NOT EFFORT. IT'S FRAGMENTATION.

The Truth About a Deal Lives in Three Different Places.



None of them talk to each other. So the CRM goes stale while the deal keeps moving.

THE INVESTIGATION BEGINS


I Went Looking for Where the Story Breaks

Real emails. Real meeting notes. Nothing invented.

BHAVIN · 30 MARCH

"All three documents attached... let me know if compliance needs anything else."

A complete, professional reply — the kind that should close a deal.



WHAT WAS REALLY GOING WRONG?

Five Whys, One Honest Answer

- 1 Deals go cold
↓
- 2 Follow-ups don't fire on time
↓
- 3 The CRM doesn't reflect reality
↓
- 4 Nobody has time to transcribe every call
↓
- 5 **Sales is founder-led — CRM upkeep is sacrificed first**

ROOT CAUSE

The CRM only captures what a busy deal owner remembers to type in.

Revenue leaks from a lack of memory — not a lack of leads.

What Kind of Tool Would Actually Help?

THE OBVIOUS TOOL

- ✗ Auto-updates the CRM
- ✗ Auto-sends follow-ups
- ✗ Black-box confidence scores
- ✗ Acts the moment it's "sure"

Fast to build. Impossible to trust.

WHAT I BUILT INSTEAD

- ✓ Read-only — never writes to the CRM
- ✓ Drafts only — a human sends
- ✓ Every claim cites a real quote
- ✓ Suggests. Never decides.

Slower to build. Something a founder can actually rely on.

Every Capability Ships With Its Guardrail

Not a checklist bolted on after — designed in from the start.

Extracts facts from raw email & note text



Grounded Attribution — must cite a verbatim quote, or it's dropped

Drafts a follow-up email, evidence-grounded



Human Control — nothing sends without a click

Summarizes the whole pipeline in seconds



Explainability — every summary traces back to source docs

Scores urgency & flags duplicates



Fairness & Privacy — business signals only, never personal data

MODEL STRATEGY

The Simplest Path That Still Holds Up

Three options. One made sense at this scale.

ELIMINATED

Fine-Tuning

No labelled data yet to do it responsibly

ELIMINATED

RAG / Embeddings

Each lead's documents are small enough to pass in full

CHOSEN

Prompt + Structured Output (Claude Sonnet 5)

Strong grounding & schema-constrained output —
nothing more was needed

When Does This Stop Being Enough?

Three specific triggers — not a guess.

TRIGGER

**A single lead outgrows
one context window**

THEN →

Add retrieval (RAG) — but only for that lead's documents, not a portfolio-wide rebuild

TRIGGER

**The question becomes
cross-lead, not per-lead**

THEN →

“Who mentioned a competitor this quarter?” needs search across leads — a dedicated insights layer, not a one-off prompt

TRIGGER

**Enough real usage data
accumulates**

THEN →

Revisit fine-tuning — once human-agreement data exists to train on responsibly

None of this is hypothetical — every trigger here is already staged in the PRD's own roadmap.

CAN THIS ACTUALLY WORK IN PRODUCTION?

The Business Case, in Real Numbers

Not a projection — this is measured, from actual billing.

\$4.85

Monthly AI cost

99.4%

Cheaper than
manual review

159.8x

Return on every
\$1 spent

\$4.85/month to run this for AtliQ's full pipeline vs. ~\$779/month in founder & rep time doing it by hand.

The cost of catching a stalled deal is smaller than the cost of one missed reply.



CLOSING THE LOOP

How We'll Know It Worked

Back to the question we opened with — one number, held up by three others.

NORTH STAR

High-value leads (>\$50k), 30+ days silent despite real engagement

TARGET → 0

OUTPUT · SHALLOW

Findings surfaced per run; suggested actions generated

Directional only — tracked, not targeted

TRUST SIGNAL

Human agreement with flagged findings; actions used as-is vs. edited vs. discarded

≥ 80% agreement

GUARDRAIL METRIC

Rate of findings dropped by the grounding check; automated CRM writes or sent emails

Writes/sends must stay at 0

One Idea, Eleven Ways It Shows Up

■ This Week's Attention

Every urgent lead, ranked and explained

■ Full Reasoning Trail

No black box — every score, explained

■ Draft Follow-Ups

One click. Nothing sends itself.

■ Effort vs. Payoff

Zero AI calls — pure CRM math

■ Drill Into a Lead

The whole story, chronologically

■ Executive Summary

The pipeline's health, in one paragraph

+ 5 more inside the live product — categories, hidden opportunities, linked evidence, memory trail, and the floating assistant.



From Question to Proof

Let's See It Live

READ-ONLY

EVIDENCE-GROUNDED

HUMAN-APPROVED, ALWAYS

ai-sales-command-center-5srn5kqnywjo2wb4dfy3u.streamlit.app