

AtliQ Sales & CRM Memory Assistant

Cost Estimation — Summary Dashboard

Monthly AI tool cost	Annual AI tool cost	Monthly savings vs. manual	% cost saved	Annual savings
\$4.85	\$58.17	\$774.55	99.4%	\$9,294.63
Analysis + summaries + emails	12 × monthly	Human-equivalent – AI cost	vs. manual review baseline	12 × monthly savings

ROI — for every \$1 spent on the AI tool, this saves:

159.8x in founder + sales rep time value

Where each number comes from

- Inference

Token pricing, per-lead/summary/email cost, monthly volume assumptions, monthly cost projection, and a low/mid/high scenario table.
- Cost Savings

Manual-review baseline, the AI tool's monthly cost (linked from Inference), the resulting savings, and a break-even calculation.
- Model & Assumptio

Why Claude Sonnet 5, infrastructure costs, the real billing data this model is grounded in, and every key assumption in plain language.

This model is fully formula-driven — every number traces back to the Inference or Cost Savings sheet; nothing here is

AtliQ Sales & CRM Memory Assistant — Cost Estimation

Model inference cost, grounded in real measured usage where possible

Yellow cells are adjustable assumptions. Blue text = input. Green text = link to another sheet.

MODEL & PRICING (Anthropic — Claude Sonnet 5)	
Model	claude-sonnet-5
Input cost (\$ / 1M tokens) — standard rate	\$3.00
Output cost (\$ / 1M tokens) — standard rate	\$15.00

Introductory rate of \$2.00 / \$10.00 per 1M tokens applies through 31 Aug 2026. The standard rate above is used throughout this model, as the more conservative, sustained estimate. Verified against Anthropic's published pricing page as of Jul 2026.

PER-LEAD ANALYSIS COST — measured from real usage, not theoretical	
Full-run cost, 43 leads, 6 categories + suggested actions (actual Anthropic billing)	\$3.00
Leads analyzed in that run	43
Average cost per lead (measured)	\$0.0698

EXECUTIVE SUMMARY COST — 1 call per Run/Refresh (estimated)	
Estimated input tokens	600
Estimated output tokens	350
Cost per summary	\$0.0071

GENERATE FOLLOW-UP EMAIL COST — on demand, per click (estimated)	
Estimated input tokens	900
Estimated output tokens	250
Cost per email generated	\$0.0065

MONTHLY VOLUME ASSUMPTIONS — steady-state AtliQ usage	
Total leads in pipeline	50
Leads needing fresh analysis / week (new email or note activity)	15
Weeks / month	4.33
Fresh lead-analyses / month	64.95
Run/Refresh clicks / month (~1 per business day)	22
Follow-up emails generated / month	25

MONTHLY COST PROJECTION	
Lead analysis cost / month	\$4.53
Executive summary cost / month	\$0.16
Follow-up email cost / month	\$0.16
Total monthly cost	\$4.85
Total annual cost	\$58.17

SCENARIO ANALYSIS — monthly cost at different fresh-analysis volumes

Scenario	Fresh leads/week	Lead analysis/mo	Exec summary/mo	Email/mo	Total/mo
Low	10	\$3.02	\$0.16	\$0.16	\$3.34
Mid (matches assumptions above)	15	\$4.53	\$0.16	\$0.16	\$4.85
High	25	\$7.55	\$0.16	\$0.16	\$7.87

Low/High scenarios vary only the fresh-analysis volume (the largest cost driver); Refresh-click and email volumes are held constant at the assumptions above.

- NOTES & ASSUMPTIONS
- 24 Jul 2026: full 43-lead run on the original 4-category schema — actual Anthropic billing: \$0.32 total.
 - Late Jul 2026, after adding Consider Deprioritizing + Suggested Next Action (6 categories): full 43-lead re-run — actual Anthropic billing: \$3.00 total. This is the figure used for the measured per-lead cost above.
 - Both figures are real, from platform.claude.com → Cost, not projections — the per-lead average already reflects actual prompt length and output verbosity, so it does not need a separate token-count estimate.
 - Executive Summary and Generate Email costs ARE estimated (not yet measured at volume), since they are newer features — token counts are conservative based on their prompt design.
 - Monthly volume assumptions describe steady-state production use by AtliQ's sales team, not this capstone's testing usage, which was much lower.
 - Caching means only leads with genuinely new email/note activity trigger a fresh API call — an unchanged lead costs \$0 on repeat runs.

Cost Savings — AI Tool vs. Manual Review

What this replaces: a founder or Jay manually re-reading emails and notes against the CRM

MANUAL EFFORT BASELINE	
Manual review time per lead needing attention (minutes)	8
Leads needing fresh review / week	15
Manual review hours / week	2.00
Weeks / month	4.33
Manual review hours / month	8.66
Blended hourly value of founder/seller time (\$)	\$90.00
Monthly human-equivalent cost	\$779.40

AI TOOL COST	
Total AI tool cost / month (linked from Inference sheet)	\$4.85

SAVINGS	
Total savings / month	\$774.55
% savings	99.4%
Annual savings	\$9,294.63

BREAK-EVEN — how much manual time would need to shrink before this stops paying off	
Break-even manual review time per lead (minutes)	0.05

At any manual review time above this many minutes/lead, the AI tool is cheaper. The assumed 8 min/lead is well above this line, so the tool is cost-positive even for very fast manual reviewers.

NOTES & ASSUMPTIONS	
• This compares the AI tool's real running cost against the time a person would spend manually cross-referencing the same leads — not against doing nothing. • 8 minutes/lead is a conservative estimate for reading linked emails/notes and comparing them to the CRM by hand; some leads (e.g. FinEdge Bank, with 3+ linked documents) would genuinely take longer. • \$90/hour is a placeholder blended rate for founder/seller time — AtliQ should replace this with a real number before relying on this figure externally. • This is a supporting, illustrative metric — the PRD's actual North Star is stale high-value leads trending to zero, not cost savings alone.	

Model Choice, Infrastructure & Key Assumptions

MODEL CHOICE	
Which model?	Claude Sonnet 5 (claude-sonnet-5), via the Anthropic API.
Why this model?	Strong instruction-following for the grounding/verbatim-quote guardrail (every finding must trace to real source text); native structured/schema-constrained output support; more than sufficient quality for extraction and comparison tasks at this scope.
Why not RAG / embeddings?	Each lead's linked-document set is small (a handful of emails/notes), so passing the full set directly in context is simpler and more reliable than building a retrieval pipeline for marginal benefit at this data volume. See AI PRD, Section 3.
Why not fine-tuning?	No labelled training data exists yet at the volume needed to fine-tune responsibly. Prompt + structured output is the right default for v1; revisit once real human-agreement data accumulates in production.
Alternate model used during prototyping?	Groq (Llama 3.3 70B) — free tier, used temporarily during development while Anthropic billing was being set up. Switched back to Claude Sonnet 5 once funded, since it's the intended production model.

INFRASTRUCTURE	
Hosting	Streamlit Community Cloud — free tier, \$0/month.
Code hosting	GitHub — free, public repository.
Storage / caching	Local JSON files (analysis_cache.json, findings_history.json) — no database cost at this scale.
Total infrastructure cost	\$0/month at prototype and early-pilot scale.

REAL MEASURED DATA USED IN THIS ESTIMATE	
24 Jul 2026	Full 43-lead run, original 4-category schema — actual Anthropic billing: \$0.32.
Late Jul 2026	Full 43-lead re-run after adding 2 categories + suggested actions (6 categories total) — actual Anthropic billing: \$3.00. Used as the basis for the measured per-lead cost in the Inference sheet.
Source	platform.claude.com → Cost — both figures are real billing data, not projections.

KEY ASSUMPTIONS	
<i>Adjustable — see yellow cells on the Inference and Cost Savings sheets.</i>	
Pipeline size	50 leads at steady state (up from the current 43, allowing for growth).
Fresh analysis volume	15 leads/week receive new email or meeting-note activity requiring re-analysis; the rest are served from cache at \$0.
Run frequency	~22 Run/Refresh clicks/month (roughly once per business day).
Email generation volume	25 draft follow-up emails generated/month, reflecting typical weekly Urgent-lead volume.